

## Problem Statement

Continuing with the same scenario, now that you have been able to successfully predict each student GPA, now you will classify each Student based on they probability to have a successful GPA score.

The different classes are:

- Low : Students where final GPA is predicted to be between: 0 and 2
- Medium : Students where final GPA is predicted to be between: 2 and 3.5
- High : Students where final GPA is predicted to be between: 3.5 and 5

## 1) Import Libraries

First let's import the following libraries, if there is any library that you need and is not in the list bellow feel free to include it

```
In [35]: # Librerías estándar
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Librerías de sklearn
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder

# Librerías de TensorFlow y Keras
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Conv1D, MaxPooling1D, Flatten
from tensorflow.keras.regularizers import l2
from tensorflow.keras.optimizers import RMSprop, Adamax
```

## 2) Load Data

- You will use the same file from the previous activity (Student Performance Data)

```
In [36]: data = pd.read_csv("M2_A3_StudentPerformanceData.csv")
```

## 3) Add a new column called 'Profile' this column will have the following information

Based on the value of GPA for each student:

- If GPA values between 0 and 2 will be labeled 'Low',
- Values between 2 and 3.5 will be 'Medium',


- And values between 3.5 and 5 will be 'High'.

```
In [37]: data['Profile'] = pd.cut(data['GPA'], bins=[0, 2, 3.5, 5], labels=['Low', 'M
```

#### 4) Use Matplotlib to show a Pie chart to show the percentage of students in each profile.

- Title: Students distribution of Profiles
- Graph Type: pie

```
In [38]: plt.figure(figsize=(10, 10))
colors = ['#a7d4ff', '#a7b7ff', '#bca7ff']
data['Profile'].value_counts().plot.pie(autopct='%1.1f%%', colors=colors)
plt.title('Profile')
plt.show()
```



## 5) Convert the Profile column into a Categorical Int

You have already created a column with three different values: 'Low', 'Medium', 'High'. These are Categorical values. But, it is important to notice that Neural Networks works better with numbers, since we apply mathematical operations to them.

Next you need to convert Profile values from Low, Medium and High, to 0, 1 and 2. IMPORTANT, the order does not matter, but make sure you always assign the same number to Low, same number to Medium and same number to High.

Make sure to use the `fit_transform` method from `LabelEncoder`.

```
In [39]: data['Profile'] = data['Profile'].replace({'Low': 0, 'Medium': 1, 'High': 2})
data = data.drop(['GPA'], axis=1)
data.dropna(inplace=True)
data['Profile'] = data['Profile'].astype(int)
```

```
<ipython-input-39-fc9e1053ecc5>:1: FutureWarning: Downcasting behavior in `r
eplace` is deprecated and will be removed in a future version. To retain the
old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in
to the future behavior, set `pd.set_option('future.no_silent_downcasting', T
rue)`
```

```
data['Profile'] = data['Profile'].replace({'Low': 0, 'Medium': 1, 'High':
2})
```

```
<ipython-input-39-fc9e1053ecc5>:1: FutureWarning: The behavior of Series.rep
lace (and DataFrame.replace) with CategoricalDtype is deprecated. In a futur
e version, replace will only be used for cases that preserve the categories.
To change the categories, use ser.cat.rename_categories instead.
```

```
data['Profile'] = data['Profile'].replace({'Low': 0, 'Medium': 1, 'High':
2})
```

```
In [40]: data
```

Out [40]:

	StudentID	Age	Gender	Ethnicity	ParentalEducation	StudyTimeWeekly	Abser
0	1001	17	1	0	2	19.833723	
1	1002	18	0	0	1	15.408756	
2	1003	15	0	2	3	4.210570	
3	1004	17	1	0	3	10.028829	
4	1005	17	1	0	2	4.672495	
...	...	...	...	...	...	...	...
2387	3388	18	1	0	3	10.680555	
2388	3389	17	0	0	1	7.583217	
2389	3390	16	1	0	2	6.805500	
2390	3391	16	1	1	0	12.416653	
2391	3392	16	1	0	2	17.819907	

2376 rows × 15 columns

## 6) Select the columns for your model.

Same as the last excersice we need a dataset for features and a dataset for label.

- Create the following dataset:
  - A dataset with the columns for the model.
  - From that data set generate the 'X' dataset. This dataset will have all the features (make sure Profile is NOT in this dataset)
  - Generate a second 'y' dataset, This dataset will only have our label column, which is 'Profile'.
  - Generate the Train and Test datasets for each X and y:
    - X\_train with 80% of the data
    - X\_test with 20% of the data
    - y\_train with 80% of the data
    - y\_test with 20% of the data

```
In [41]: X = data.drop(['Profile'], axis=1)
y = data['Profile']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran
```

## 7) All Feature datasets in the same scale.

Use StandardScaler to make sure all features in the X\_train and X\_test datasets are on the same scale.

Standardization transforms your data so that it has a mean of 0 and a standard deviation of 1. This is important because many machine learning algorithms perform better when the input features are on a similar scale.

Reason for Using StandardScaler:

- Consistent Scale: Features with different scales (e.g., age in years, income in dollars) can bias the model. StandardScaler ensures all features contribute equally.
- Improved Convergence: Algorithms like gradient descent converge faster with standardized data.
- Regularization: Helps in achieving better performance in regularization methods like Ridge and Lasso regression.

```
In [42]: scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

## 8. Define your Deep Neural Network.

- This will be a Sequential Neural Network.
- With a Dense input layer with 64 units, and input dimension based on the X\_train size and Relu as the activation function.
- A Dense hidden layer with 32 units, and Relu as the activation function.
- And a Dense output layer with the number of different values in the y dataset, activation function = to softmax

This last part of the output layer is super important, since we want to do a classification and not a regression, we will use activation functions that fits better a classification scenario.

```
In [43]: model = Sequential()
model.add(Dense(64, input_dim=X_train.shape[1], activation='relu'))
model.add(Dense(32, activation='relu'))
model.add(Dense(y_train.unique(), activation='softmax'))
```

```
/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
```

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

## 9. Compile your Neural Network


- Choose Adam as the optimizer
- And sparse\_categorical\_crossentropy as the Loss function
- Also add the following metrics: accuracy


```
In [44]: model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metr
```


## 10. Fit (or train) your model


- Use the X\_train and y\_train datasets for the training
- Do 50 data iterations
- Choose the batch size = 10
- Also select a validation\_split of 0.2
- Save the result of the fit function in a variable called 'history'


```
In [45]: history = model.fit(X_train, y_train, epochs=50, batch_size=10, validation_s
```


Epoch 1/50  
152/152  2s 3ms/step - accuracy: 0.6627 - loss: 0.7972 -  
val\_accuracy: 0.8816 - val\_loss: 0.3778


Epoch 2/50  
152/152  0s 2ms/step - accuracy: 0.9189 - loss: 0.2861 -  
val\_accuracy: 0.9237 - val\_loss: 0.2406


Epoch 3/50  
152/152  0s 2ms/step - accuracy: 0.9476 - loss: 0.1636 -  
val\_accuracy: 0.9289 - val\_loss: 0.1918


Epoch 4/50  
152/152  1s 2ms/step - accuracy: 0.9592 - loss: 0.1247 -  
val\_accuracy: 0.9474 - val\_loss: 0.1635


Epoch 5/50  
152/152  1s 2ms/step - accuracy: 0.9772 - loss: 0.0858 -  
val\_accuracy: 0.9526 - val\_loss: 0.1431


Epoch 6/50  
152/152  0s 2ms/step - accuracy: 0.9839 - loss: 0.0735 -  
val\_accuracy: 0.9474 - val\_loss: 0.1399


Epoch 7/50  
152/152  0s 2ms/step - accuracy: 0.9809 - loss: 0.0654 -  
val\_accuracy: 0.9553 - val\_loss: 0.1301


Epoch 8/50  
152/152  0s 2ms/step - accuracy: 0.9857 - loss: 0.0648 -  
val\_accuracy: 0.9579 - val\_loss: 0.1232


Epoch 9/50  
152/152  1s 3ms/step - accuracy: 0.9866 - loss: 0.0402 -  
val\_accuracy: 0.9579 - val\_loss: 0.1215


Epoch 10/50  
152/152  1s 3ms/step - accuracy: 0.9927 - loss: 0.0337 -  
val\_accuracy: 0.9605 - val\_loss: 0.1242


Epoch 11/50  
152/152  1s 3ms/step - accuracy: 0.9900 - loss: 0.0362 -  
val\_accuracy: 0.9632 - val\_loss: 0.1156


Epoch 12/50  
152/152  0s 3ms/step - accuracy: 0.9944 - loss: 0.0255 -  
val\_accuracy: 0.9632 - val\_loss: 0.1163


Epoch 13/50  
152/152  0s 2ms/step - accuracy: 0.9965 - loss: 0.0223 -  
val\_accuracy: 0.9632 - val\_loss: 0.1172


Epoch 14/50  
152/152  0s 2ms/step - accuracy: 0.9960 - loss: 0.0181 -  
val\_accuracy: 0.9579 - val\_loss: 0.1269



















Epoch 15/50  
152/152  0s 2ms/step - accuracy: 0.9974 - loss: 0.0159 -  
val\_accuracy: 0.9658 - val\_loss: 0.1216

Epoch 16/50  
152/152  0s 2ms/step - accuracy: 0.9985 - loss: 0.0145 -  
val\_accuracy: 0.9658 - val\_loss: 0.1211


Epoch 17/50  
152/152  0s 2ms/step - accuracy: 0.9993 - loss: 0.0103 -  
val\_accuracy: 0.9605 - val\_loss: 0.1236


Epoch 18/50  
152/152  1s 2ms/step - accuracy: 0.9998 - loss: 0.0071 -  
val\_accuracy: 0.9605 - val\_loss: 0.1306


Epoch 19/50  
152/152  1s 2ms/step - accuracy: 0.9999 - loss: 0.0063 -


```
val_accuracy: 0.9605 - val_loss: 0.1260
Epoch 20/50
152/152  0s 2ms/step - accuracy: 0.9985 - loss: 0.0068 -
val_accuracy: 0.9684 - val_loss: 0.1334
Epoch 21/50
152/152  1s 2ms/step - accuracy: 1.0000 - loss: 0.0067 -
val_accuracy: 0.9605 - val_loss: 0.1379
Epoch 22/50
152/152  0s 2ms/step - accuracy: 1.0000 - loss: 0.0045 -
val_accuracy: 0.9605 - val_loss: 0.1395
Epoch 23/50
152/152  0s 2ms/step - accuracy: 1.0000 - loss: 0.0047 -
val_accuracy: 0.9632 - val_loss: 0.1421
Epoch 24/50
152/152  1s 2ms/step - accuracy: 1.0000 - loss: 0.0031 -
val_accuracy: 0.9605 - val_loss: 0.1359
Epoch 25/50
152/152  0s 2ms/step - accuracy: 1.0000 - loss: 0.0030 -
val_accuracy: 0.9658 - val_loss: 0.1427
Epoch 26/50
152/152  1s 2ms/step - accuracy: 1.0000 - loss: 0.0027 -
val_accuracy: 0.9632 - val_loss: 0.1526
Epoch 27/50
152/152  0s 2ms/step - accuracy: 1.0000 - loss: 0.0020 -
val_accuracy: 0.9632 - val_loss: 0.1423
Epoch 28/50
152/152  0s 2ms/step - accuracy: 1.0000 - loss: 0.0018 -
val_accuracy: 0.9658 - val_loss: 0.1500
Epoch 29/50
152/152  1s 2ms/step - accuracy: 1.0000 - loss: 0.0016 -
val_accuracy: 0.9632 - val_loss: 0.1548
Epoch 30/50
152/152  0s 2ms/step - accuracy: 1.0000 - loss: 0.0014 -
val_accuracy: 0.9658 - val_loss: 0.1477
Epoch 31/50
152/152  1s 2ms/step - accuracy: 1.0000 - loss: 0.0011 -
val_accuracy: 0.9632 - val_loss: 0.1603
Epoch 32/50
152/152  0s 2ms/step - accuracy: 1.0000 - loss: 0.0011 -
val_accuracy: 0.9658 - val_loss: 0.1581
Epoch 33/50
152/152  0s 2ms/step - accuracy: 1.0000 - loss: 8.5025
e-04 - val_accuracy: 0.9605 - val_loss: 0.1691
Epoch 34/50
152/152  0s 2ms/step - accuracy: 1.0000 - loss: 8.4901
e-04 - val_accuracy: 0.9658 - val_loss: 0.1664
Epoch 35/50
152/152  0s 2ms/step - accuracy: 1.0000 - loss: 6.5848
e-04 - val_accuracy: 0.9632 - val_loss: 0.1742
Epoch 36/50
152/152  0s 2ms/step - accuracy: 1.0000 - loss: 5.5531
e-04 - val_accuracy: 0.9632 - val_loss: 0.1725
Epoch 37/50
152/152  0s 2ms/step - accuracy: 1.0000 - loss: 5.1014
e-04 - val_accuracy: 0.9632 - val_loss: 0.1727
Epoch 38/50
```





152/152  0s 3ms/step - accuracy: 1.0000 - loss: 4.1421  
e-04 - val\_accuracy: 0.9632 - val\_loss: 0.1776  
Epoch 39/50


152/152  1s 3ms/step - accuracy: 1.0000 - loss: 4.2511  
e-04 - val\_accuracy: 0.9632 - val\_loss: 0.1844  
Epoch 40/50


152/152  1s 3ms/step - accuracy: 1.0000 - loss: 3.9371  
e-04 - val\_accuracy: 0.9632 - val\_loss: 0.1872  
Epoch 41/50


152/152  1s 3ms/step - accuracy: 1.0000 - loss: 3.0134  
e-04 - val\_accuracy: 0.9658 - val\_loss: 0.1791  
Epoch 42/50


152/152  1s 3ms/step - accuracy: 1.0000 - loss: 2.5413  
e-04 - val\_accuracy: 0.9632 - val\_loss: 0.1892  
Epoch 43/50


152/152  0s 2ms/step - accuracy: 1.0000 - loss: 2.4022  
e-04 - val\_accuracy: 0.9658 - val\_loss: 0.1869  
Epoch 44/50


152/152  0s 2ms/step - accuracy: 1.0000 - loss: 2.3124  
e-04 - val\_accuracy: 0.9632 - val\_loss: 0.1889  
Epoch 45/50


152/152  0s 2ms/step - accuracy: 1.0000 - loss: 2.3569  
e-04 - val\_accuracy: 0.9632 - val\_loss: 0.1990  
Epoch 46/50

152/152  0s 2ms/step - accuracy: 1.0000 - loss: 1.8897  
e-04 - val\_accuracy: 0.9632 - val\_loss: 0.1929  
Epoch 47/50

152/152  0s 2ms/step - accuracy: 1.0000 - loss: 1.8736  
e-04 - val\_accuracy: 0.9605 - val\_loss: 0.1968  
Epoch 48/50

152/152  1s 2ms/step - accuracy: 1.0000 - loss: 1.4987  
e-04 - val\_accuracy: 0.9632 - val\_loss: 0.2025  
Epoch 49/50


152/152  0s 2ms/step - accuracy: 1.0000 - loss: 1.5218  
e-04 - val\_accuracy: 0.9632 - val\_loss: 0.1970  
Epoch 50/50

152/152  0s 2ms/step - accuracy: 1.0000 - loss: 1.3149  
e-04 - val\_accuracy: 0.9605 - val\_loss: 0.2010

## 11. View your history variable:

- Use Matplotlib.pyplot to show graphs of your model training history
- In one graph:
  - Plot the Training Accuracy and the Validation Accuracy
  - X Label = Epochs
  - Y Label = Accuracy
  - Title = Model Accuracy over Epochs
- In a second graph:
  - Plot the Training Loss and the Validation Loss
  - X Label = Epochs
  - Y Label = Loss
  - Title = Model Loss over Epochs

```
In [46]: plt.plot(history.history['accuracy'], color='#a7c9ff')
plt.plot(history.history['val_accuracy'], color='#bca7ff')
plt.title('Model Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```



## 12. Evaluate your model:

- See the result of your loss function.
- What can you deduct from there?

The model stops learning after 30 epochs because the loss function no longer shows a decrease in the training and validation data sets. After that point, the value of the loss remains nearly constant, indicating that the model has reached a point where it cannot improve its performance, possibly because it has reached a state of convergence or a local minimum. This behavior suggests that further training will not result in significant improvements in model accuracy or generalization.

```
In [47]: accuracy = model.evaluate(X_test, y_test)
print(f"Accuracy: {accuracy[1]}")
```

15/15 ————— 0s 1ms/step – accuracy: 0.9436 – loss: 0.3663  
Accuracy: 0.9516806602478027

The model has achieved satisfactory performance with an accuracy of 95.17% in the test set.

However, because the loss function stops decreasing after 30 epochs in both the training and validation data, the model has probably reached a saturation point in its learning. This indicates that further training will not result in further improvements in the accuracy or generalization ability of the model, and could even lead to overfitting. Therefore, stopping training at this point is appropriate to avoid diminishing returns.

### 13. Use your model to make some predictions:

- Make predictions of your X\_test dataset
- Print the each of the predictions and the actual value (which is in y\_test)
- Replace the 'Low', 'Medium' and 'High' to your actual and predicted values.
- How good was your model?

```
In [48]: y_pred = model.predict(X_test)
```

15/15 ————— 0s 5ms/step

```
In [49]: # Create a DataFrame for the predicted probabilities and determine the predicted class
predicted_probabilities = pd.DataFrame(y_pred, columns=['Low', 'Medium', 'High'])
predicted_probabilities['Predicted_Class'] = predicted_probabilities.idxmax(axis=1)

# Add the actual class labels from the test set
predicted_probabilities['Actual_Class'] = y_test.values

# Display the DataFrame with predictions and actual values
predicted_probabilities.head() # Use .head() to display the first few rows
```

Out[49]:

		Low	Medium	High	Predicted_Class	Actual_Class
0	2.954267e-07	9.999240e-01	7.570768e-05		Medium	1
1	6.375265e-09	9.999999e-01	4.693620e-10		Medium	1
2	3.242281e-10	9.999992e-01	7.066138e-07		Medium	1
3	9.999999e-01	7.275351e-17	4.602823e-20		Low	0
4	9.999999e-01	1.633048e-15	5.726148e-18		Low	0

## 14. Compete against this model:

- Create two more different models to compete with this model
- Here are a few ideas of things you can change:
  - During Dataset data engineering:
    - You can remove features that you think do not help in the training and prediction
    - Feature Scaling: Ensure all features are on a similar scale (as you already did with StandardScaler)
  - During Model Definition:
    - You can change the Model Architecture (change the type or number of layers or the number of units)
    - You can add dropout layers to prevent overfitting
  - During Model Compile:
    - You can try other optimizer when compiling your model, here some optimizer samples: Adam, RMSprop, or Adagrad.
    - Try another Loss Function
  - During Model Training:
    - Encrease the number of Epochs
    - Adjust the size of your batch
- Explain in a Markdown cell which changes are you implementing
- Show the comparison of your model versus the original model

### Model 2:

- Changes:
  - Dataset Data Engineering
  - Model Definition
  - Model Compile
  - Model Training

Changes in the architecture of the model 2:

- The model has 2 hidden layers with 1 dense with 32 with 2 dropouts after each dense

layer.

- The third hidden layer has just 16 neurons.
- The rest of the process is the same as the original model.

```
In [50]: model_2 = Sequential()
model_2.add(Dense(64, input_dim=X_train.shape[1], activation='relu'))
model_2.add(Dropout(0.4))
model_2.add(Dense(64, activation='relu'))
model_2.add(Dropout(0.4))
model_2.add(Dense(32, activation='relu'))
model_2.add(Dense(16, activation='relu'))
model_2.add(Dense(y_train.nunique(), activation='softmax'))
optimizer = RMSprop(learning_rate=0.001)

model_2.compile(optimizer=optimizer, loss='sparse_categorical_crossentropy',


history_2 = model_2.fit(X_train, y_train, epochs=50, batch_size=10, validati


# show the history
plt.plot(history_2.history['accuracy'], color='#a7c9ff')
plt.plot(history_2.history['val_accuracy'], color='#bca7ff')
plt.title('Model Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```


Epoch 1/50


/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.


```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```


152/152  1s 3ms/step - accuracy: 0.6751 - loss: 0.7515 -  
val\_accuracy: 0.8737 - val\_loss: 0.3690  
Epoch 2/50


152/152  0s 2ms/step - accuracy: 0.8567 - loss: 0.4304 -  
val\_accuracy: 0.9105 - val\_loss: 0.2882  
Epoch 3/50


152/152  1s 2ms/step - accuracy: 0.8933 - loss: 0.3327 -  
val\_accuracy: 0.9079 - val\_loss: 0.2773  
Epoch 4/50


152/152  1s 2ms/step - accuracy: 0.9058 - loss: 0.2880 -  
val\_accuracy: 0.9289 - val\_loss: 0.2299  
Epoch 5/50


152/152  1s 2ms/step - accuracy: 0.9261 - loss: 0.2431 -  
val\_accuracy: 0.9263 - val\_loss: 0.2261  
Epoch 6/50


152/152  1s 2ms/step - accuracy: 0.9163 - loss: 0.2433 -  
val\_accuracy: 0.9316 - val\_loss: 0.2110  
Epoch 7/50


152/152  0s 2ms/step - accuracy: 0.9328 - loss: 0.2240 -  
val\_accuracy: 0.9237 - val\_loss: 0.2045  
Epoch 8/50


152/152  0s 2ms/step - accuracy: 0.9278 - loss: 0.1995 -  
val\_accuracy: 0.9368 - val\_loss: 0.1845  
Epoch 9/50

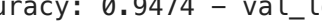
152/152  0s 2ms/step - accuracy: 0.9527 - loss: 0.1679 -  
val\_accuracy: 0.9395 - val\_loss: 0.1765  
Epoch 10/50


152/152  1s 2ms/step - accuracy: 0.9440 - loss: 0.1715 -  
val\_accuracy: 0.9395 - val\_loss: 0.1619  
Epoch 11/50

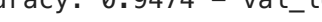
152/152  0s 2ms/step - accuracy: 0.9399 - loss: 0.1780 -  
val\_accuracy: 0.9447 - val\_loss: 0.1562  
Epoch 12/50


152/152  0s 3ms/step - accuracy: 0.9287 - loss: 0.2375 -  
val\_accuracy: 0.9395 - val\_loss: 0.1869  
Epoch 13/50


152/152  0s 3ms/step - accuracy: 0.9475 - loss: 0.1619 -  
val\_accuracy: 0.9447 - val\_loss: 0.1527  
Epoch 14/50


152/152  1s 3ms/step - accuracy: 0.9619 - loss: 0.1186 -  
val\_accuracy: 0.9474 - val\_loss: 0.1749  
Epoch 15/50


152/152  0s 3ms/step - accuracy: 0.9537 - loss: 0.1344 -  
val\_accuracy: 0.9553 - val\_loss: 0.1582  
Epoch 16/50


152/152  1s 4ms/step - accuracy: 0.9573 - loss: 0.1658 -  
val\_accuracy: 0.9474 - val\_loss: 0.1635  
Epoch 17/50


152/152  0s 2ms/step - accuracy: 0.9474 - loss: 0.1698 -  
val\_accuracy: 0.9500 - val\_loss: 0.1546  
Epoch 18/50


152/152  0s 2ms/step - accuracy: 0.9640 - loss: 0.1139 -  
val\_accuracy: 0.9447 - val\_loss: 0.1588  
Epoch 19/50


152/152  1s 2ms/step - accuracy: 0.9546 - loss: 0.1385 -  
val\_accuracy: 0.9421 - val\_loss: 0.1804


Epoch 20/50  
152/152  0s 2ms/step - accuracy: 0.9550 - loss: 0.1425 -  
val\_accuracy: 0.9500 - val\_loss: 0.1484


Epoch 21/50  
152/152  1s 2ms/step - accuracy: 0.9542 - loss: 0.1370 -  
val\_accuracy: 0.9447 - val\_loss: 0.1860


Epoch 22/50  
152/152  1s 2ms/step - accuracy: 0.9664 - loss: 0.1101 -  
val\_accuracy: 0.9421 - val\_loss: 0.1701


Epoch 23/50  
152/152  1s 2ms/step - accuracy: 0.9672 - loss: 0.1089 -  
val\_accuracy: 0.9447 - val\_loss: 0.1537


Epoch 24/50  
152/152  0s 2ms/step - accuracy: 0.9641 - loss: 0.1112 -  
val\_accuracy: 0.9316 - val\_loss: 0.2378


Epoch 25/50  
152/152  0s 2ms/step - accuracy: 0.9525 - loss: 0.1391 -  
val\_accuracy: 0.9500 - val\_loss: 0.1703


Epoch 26/50  
152/152  0s 2ms/step - accuracy: 0.9658 - loss: 0.1105 -  
val\_accuracy: 0.9526 - val\_loss: 0.1597


Epoch 27/50  
152/152  0s 2ms/step - accuracy: 0.9644 - loss: 0.0998 -  
val\_accuracy: 0.9474 - val\_loss: 0.1804


Epoch 28/50  
152/152  0s 2ms/step - accuracy: 0.9647 - loss: 0.1177 -  
val\_accuracy: 0.9447 - val\_loss: 0.1590


Epoch 29/50  
152/152  0s 2ms/step - accuracy: 0.9635 - loss: 0.1327 -  
val\_accuracy: 0.9474 - val\_loss: 0.1532


Epoch 30/50  
152/152  0s 2ms/step - accuracy: 0.9767 - loss: 0.0851 -  
val\_accuracy: 0.9421 - val\_loss: 0.1871


Epoch 31/50  
152/152  1s 2ms/step - accuracy: 0.9652 - loss: 0.1156 -  
val\_accuracy: 0.9474 - val\_loss: 0.1828


Epoch 32/50  
152/152  1s 2ms/step - accuracy: 0.9712 - loss: 0.1025 -  
val\_accuracy: 0.9447 - val\_loss: 0.1777


Epoch 33/50  
152/152  0s 2ms/step - accuracy: 0.9649 - loss: 0.1162 -  
val\_accuracy: 0.9579 - val\_loss: 0.1479













Epoch 34/50  
152/152  0s 2ms/step - accuracy: 0.9638 - loss: 0.1088 -  
val\_accuracy: 0.9526 - val\_loss: 0.1592

Epoch 35/50  
152/152  0s 2ms/step - accuracy: 0.9660 - loss: 0.1019 -  
val\_accuracy: 0.9500 - val\_loss: 0.1739


Epoch 36/50  
152/152  0s 2ms/step - accuracy: 0.9656 - loss: 0.0938 -  
val\_accuracy: 0.9526 - val\_loss: 0.1604

Epoch 37/50  
152/152  1s 2ms/step - accuracy: 0.9591 - loss: 0.1155 -  
val\_accuracy: 0.9553 - val\_loss: 0.1496

Epoch 38/50  
152/152  0s 2ms/step - accuracy: 0.9547 - loss: 0.1316 -

val\_accuracy: 0.9447 - val\_loss: 0.1650  
Epoch 39/50  
152/152  0s 2ms/step - accuracy: 0.9628 - loss: 0.1171 -  
val\_accuracy: 0.9553 - val\_loss: 0.1438  
Epoch 40/50  
152/152  1s 2ms/step - accuracy: 0.9720 - loss: 0.0817 -  
val\_accuracy: 0.9500 - val\_loss: 0.1988  
Epoch 41/50  
152/152  0s 3ms/step - accuracy: 0.9786 - loss: 0.0782 -  
val\_accuracy: 0.9447 - val\_loss: 0.2135  
Epoch 42/50  
152/152  1s 3ms/step - accuracy: 0.9691 - loss: 0.1314 -  
val\_accuracy: 0.9526 - val\_loss: 0.1557  
Epoch 43/50  
152/152  1s 3ms/step - accuracy: 0.9731 - loss: 0.1107 -  
val\_accuracy: 0.9579 - val\_loss: 0.1507  
Epoch 44/50  
152/152  1s 3ms/step - accuracy: 0.9835 - loss: 0.0652 -  
val\_accuracy: 0.9447 - val\_loss: 0.2514  
Epoch 45/50  
152/152  1s 4ms/step - accuracy: 0.9719 - loss: 0.0913 -  
val\_accuracy: 0.9526 - val\_loss: 0.1722  
Epoch 46/50  
152/152  0s 2ms/step - accuracy: 0.9683 - loss: 0.1083 -  
val\_accuracy: 0.9500 - val\_loss: 0.1607  
Epoch 47/50  
152/152  1s 2ms/step - accuracy: 0.9702 - loss: 0.1013 -  
val\_accuracy: 0.9553 - val\_loss: 0.1524  
Epoch 48/50  
152/152  0s 2ms/step - accuracy: 0.9733 - loss: 0.0910 -  
val\_accuracy: 0.9526 - val\_loss: 0.1541  
Epoch 49/50  
152/152  0s 2ms/step - accuracy: 0.9813 - loss: 0.0764 -  
val\_accuracy: 0.9526 - val\_loss: 0.1769  
Epoch 50/50  
152/152  1s 2ms/step - accuracy: 0.9626 - loss: 0.1270 -  
val\_accuracy: 0.9553 - val\_loss: 0.1587





```
In [51]: accuracy_2 = model_2.evaluate(X_test, y_test)
print(f"Accuracy: {accuracy_2[1]}")
```

15/15 ————— 0s 1ms/step – accuracy: 0.9390 – loss: 0.1700  
Accuracy: 0.9495798349380493

### Model 3:

- Changes:
  - Dataset Data Engineering
  - Model Definition
  - Model Compile
  - Model Training

Changes in the architecture of the model 3:

- The model has 4 hidden layers with 2 dense with 32.
- One of the hidden layers has 16 neurons and the other has 64 neurons.
- Increased the number of epochs to 100.
- The rest of the process is the same as the original model.

```
In [52]: model_3 = Sequential()
model_3.add(Dense(32, input_dim=X_train.shape[1], activation='relu'))
model_3.add(Dense(64, activation='relu'))
model_3.add(Dropout(0.3))
```


```
model_3.add(Dense(128, activation='relu'))
model_3.add(Dropout(0.3))
model_3.add(Dense(64, activation='relu'))
model_3.add(Dense(32, activation='relu'))
model_3.add(Dense(y_train.nunique(), activation='softmax'))


optimizer = Adamax(learning_rate=0.002)


model_3.compile(optimizer=optimizer, loss='sparse_categorical_crossentropy',


history_3 = model_3.fit(X_train, y_train, epochs=100, batch_size=10, validat


# show the history
plt.plot(history_3.history['accuracy'], color='#a7c9ff')
plt.plot(history_3.history['val_accuracy'], color='#bca7ff')
plt.title('Model Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```


Epoch 1/100  
152/152  2s 3ms/step - accuracy: 0.6834 - loss: 0.7580 -  
val\_accuracy: 0.8868 - val\_loss: 0.3300


Epoch 2/100  
152/152  0s 2ms/step - accuracy: 0.8937 - loss: 0.3162 -  
val\_accuracy: 0.9026 - val\_loss: 0.2710


Epoch 3/100  
152/152  0s 2ms/step - accuracy: 0.9221 - loss: 0.2421 -  
val\_accuracy: 0.9158 - val\_loss: 0.2343


Epoch 4/100  
152/152  0s 2ms/step - accuracy: 0.9268 - loss: 0.2258 -  
val\_accuracy: 0.9184 - val\_loss: 0.2110


Epoch 5/100  
152/152  1s 2ms/step - accuracy: 0.9385 - loss: 0.2086 -  
val\_accuracy: 0.9211 - val\_loss: 0.2031


Epoch 6/100  
152/152  0s 2ms/step - accuracy: 0.9386 - loss: 0.1669 -  
val\_accuracy: 0.9237 - val\_loss: 0.1849


Epoch 7/100  
152/152  1s 2ms/step - accuracy: 0.9447 - loss: 0.1645 -  
val\_accuracy: 0.9447 - val\_loss: 0.1682


Epoch 8/100  
152/152  0s 2ms/step - accuracy: 0.9355 - loss: 0.1621 -  
val\_accuracy: 0.9316 - val\_loss: 0.1642


Epoch 9/100  
152/152  1s 2ms/step - accuracy: 0.9535 - loss: 0.1397 -  
val\_accuracy: 0.9421 - val\_loss: 0.1671


Epoch 10/100  
152/152  0s 2ms/step - accuracy: 0.9534 - loss: 0.1279 -  
val\_accuracy: 0.9447 - val\_loss: 0.1538


Epoch 11/100  
152/152  0s 2ms/step - accuracy: 0.9592 - loss: 0.1213 -  
val\_accuracy: 0.9474 - val\_loss: 0.1510


Epoch 12/100  
152/152  0s 2ms/step - accuracy: 0.9620 - loss: 0.1061 -  
val\_accuracy: 0.9474 - val\_loss: 0.1441


Epoch 13/100  
152/152  0s 2ms/step - accuracy: 0.9513 - loss: 0.1325 -  
val\_accuracy: 0.9500 - val\_loss: 0.1402


Epoch 14/100  
152/152  1s 3ms/step - accuracy: 0.9631 - loss: 0.0957 -  
val\_accuracy: 0.9553 - val\_loss: 0.1411



















Epoch 15/100  
152/152  1s 4ms/step - accuracy: 0.9629 - loss: 0.1188 -  
val\_accuracy: 0.9421 - val\_loss: 0.1384


Epoch 16/100  
152/152  1s 3ms/step - accuracy: 0.9534 - loss: 0.1278 -  
val\_accuracy: 0.9421 - val\_loss: 0.1746


Epoch 17/100  
152/152  1s 4ms/step - accuracy: 0.9668 - loss: 0.1006 -  
val\_accuracy: 0.9447 - val\_loss: 0.1377


Epoch 18/100  
152/152  1s 2ms/step - accuracy: 0.9715 - loss: 0.0878 -  
val\_accuracy: 0.9526 - val\_loss: 0.1417


Epoch 19/100  
152/152  1s 2ms/step - accuracy: 0.9657 - loss: 0.0912 -


```
val_accuracy: 0.9500 - val_loss: 0.1511
Epoch 20/100
152/152  0s 2ms/step - accuracy: 0.9665 - loss: 0.0864 -
val_accuracy: 0.9553 - val_loss: 0.1449
Epoch 21/100
152/152  0s 2ms/step - accuracy: 0.9658 - loss: 0.0989 -
val_accuracy: 0.9474 - val_loss: 0.1365
Epoch 22/100
152/152  0s 2ms/step - accuracy: 0.9606 - loss: 0.0911 -
val_accuracy: 0.9474 - val_loss: 0.1431
Epoch 23/100
152/152  1s 2ms/step - accuracy: 0.9685 - loss: 0.0806 -
val_accuracy: 0.9500 - val_loss: 0.1287
Epoch 24/100
152/152  0s 2ms/step - accuracy: 0.9736 - loss: 0.0693 -
val_accuracy: 0.9605 - val_loss: 0.1298
Epoch 25/100
152/152  1s 2ms/step - accuracy: 0.9810 - loss: 0.0495 -
val_accuracy: 0.9632 - val_loss: 0.1367
Epoch 26/100
152/152  1s 2ms/step - accuracy: 0.9775 - loss: 0.0665 -
val_accuracy: 0.9553 - val_loss: 0.1253
Epoch 27/100
152/152  0s 2ms/step - accuracy: 0.9690 - loss: 0.0878 -
val_accuracy: 0.9421 - val_loss: 0.1552
Epoch 28/100
152/152  1s 2ms/step - accuracy: 0.9742 - loss: 0.0593 -
val_accuracy: 0.9500 - val_loss: 0.1479
Epoch 29/100
152/152  0s 2ms/step - accuracy: 0.9751 - loss: 0.0597 -
val_accuracy: 0.9579 - val_loss: 0.1319
Epoch 30/100
152/152  0s 2ms/step - accuracy: 0.9839 - loss: 0.0440 -
val_accuracy: 0.9474 - val_loss: 0.1459
Epoch 31/100
152/152  1s 2ms/step - accuracy: 0.9714 - loss: 0.0701 -
val_accuracy: 0.9553 - val_loss: 0.1288
Epoch 32/100
152/152  0s 2ms/step - accuracy: 0.9822 - loss: 0.0610 -
val_accuracy: 0.9526 - val_loss: 0.1285
Epoch 33/100
152/152  1s 3ms/step - accuracy: 0.9802 - loss: 0.0673 -
val_accuracy: 0.9632 - val_loss: 0.1220
Epoch 34/100
152/152  1s 3ms/step - accuracy: 0.9884 - loss: 0.0377 -
val_accuracy: 0.9500 - val_loss: 0.1420
Epoch 35/100
152/152  0s 2ms/step - accuracy: 0.9806 - loss: 0.0623 -
val_accuracy: 0.9526 - val_loss: 0.1371
Epoch 36/100
152/152  1s 2ms/step - accuracy: 0.9768 - loss: 0.0628 -
val_accuracy: 0.9632 - val_loss: 0.1164
Epoch 37/100
152/152  1s 4ms/step - accuracy: 0.9923 - loss: 0.0284 -
val_accuracy: 0.9684 - val_loss: 0.1278
Epoch 38/100
```


152/152  1s 4ms/step - accuracy: 0.9828 - loss: 0.0480 -  
val\_accuracy: 0.9526 - val\_loss: 0.1395  
Epoch 39/100


152/152  1s 3ms/step - accuracy: 0.9818 - loss: 0.0566 -  
val\_accuracy: 0.9421 - val\_loss: 0.1729  
Epoch 40/100


152/152  1s 3ms/step - accuracy: 0.9794 - loss: 0.0572 -  
val\_accuracy: 0.9553 - val\_loss: 0.1448  
Epoch 41/100


152/152  1s 3ms/step - accuracy: 0.9834 - loss: 0.0399 -  
val\_accuracy: 0.9553 - val\_loss: 0.1338  
Epoch 42/100


152/152  0s 2ms/step - accuracy: 0.9826 - loss: 0.0427 -  
val\_accuracy: 0.9553 - val\_loss: 0.1361  
Epoch 43/100


152/152  1s 2ms/step - accuracy: 0.9847 - loss: 0.0429 -  
val\_accuracy: 0.9553 - val\_loss: 0.1511  
Epoch 44/100

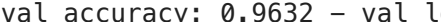
152/152  1s 2ms/step - accuracy: 0.9850 - loss: 0.0418 -  
val\_accuracy: 0.9579 - val\_loss: 0.1505  
Epoch 45/100

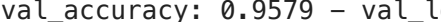
152/152  1s 2ms/step - accuracy: 0.9841 - loss: 0.0451 -  
val\_accuracy: 0.9684 - val\_loss: 0.1198  
Epoch 46/100

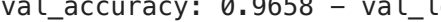
152/152  1s 2ms/step - accuracy: 0.9736 - loss: 0.0560 -  
val\_accuracy: 0.9605 - val\_loss: 0.1302  
Epoch 47/100

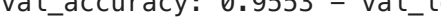
152/152  1s 2ms/step - accuracy: 0.9864 - loss: 0.0407 -  
val\_accuracy: 0.9553 - val\_loss: 0.1361  
Epoch 48/100

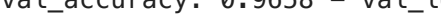
152/152  0s 2ms/step - accuracy: 0.9809 - loss: 0.0497 -  
val\_accuracy: 0.9605 - val\_loss: 0.1193  
Epoch 49/100


152/152  0s 2ms/step - accuracy: 0.9845 - loss: 0.0384 -  
val\_accuracy: 0.9632 - val\_loss: 0.1271  
Epoch 50/100


152/152  0s 2ms/step - accuracy: 0.9921 - loss: 0.0287 -  
val\_accuracy: 0.9579 - val\_loss: 0.1242  
Epoch 51/100


152/152  0s 2ms/step - accuracy: 0.9854 - loss: 0.0343 -  
val\_accuracy: 0.9658 - val\_loss: 0.1137  
Epoch 52/100

152/152  0s 2ms/step - accuracy: 0.9898 - loss: 0.0337 -  
val\_accuracy: 0.9553 - val\_loss: 0.1551  
Epoch 53/100


152/152  0s 2ms/step - accuracy: 0.9880 - loss: 0.0344 -  
val\_accuracy: 0.9658 - val\_loss: 0.1185  
Epoch 54/100

152/152  0s 2ms/step - accuracy: 0.9904 - loss: 0.0257 -  
val\_accuracy: 0.9711 - val\_loss: 0.1228  
Epoch 55/100

152/152  1s 2ms/step - accuracy: 0.9873 - loss: 0.0311 -  
val\_accuracy: 0.9605 - val\_loss: 0.1333  
Epoch 56/100

152/152  0s 2ms/step - accuracy: 0.9879 - loss: 0.0286 -  
val\_accuracy: 0.9632 - val\_loss: 0.1293


Epoch 57/100

**152/152**  **1s** 2ms/step - accuracy: 0.9860 - loss: 0.0323 - val\_accuracy: 0.9605 - val\_loss: 0.1513

Epoch 58/100

**152/152**  **1s** 2ms/step - accuracy: 0.9893 - loss: 0.0338 - val\_accuracy: 0.9632 - val\_loss: 0.1295

Epoch 59/100

**152/152**  **0s** 2ms/step - accuracy: 0.9939 - loss: 0.0140 - val\_accuracy: 0.9605 - val\_loss: 0.1571


Epoch 60/100

**152/152**  **0s** 2ms/step - accuracy: 0.9846 - loss: 0.0356 - val\_accuracy: 0.9605 - val\_loss: 0.1485


Epoch 61/100

**152/152**  **1s** 2ms/step - accuracy: 0.9913 - loss: 0.0319 - val\_accuracy: 0.9605 - val\_loss: 0.1386


Epoch 62/100

**152/152**  **1s** 3ms/step - accuracy: 0.9936 - loss: 0.0160 - val\_accuracy: 0.9632 - val\_loss: 0.1409


Epoch 63/100

**152/152**  **1s** 3ms/step - accuracy: 0.9876 - loss: 0.0282 - val\_accuracy: 0.9632 - val\_loss: 0.1324


Epoch 64/100

**152/152**  **1s** 4ms/step - accuracy: 0.9876 - loss: 0.0242 - val\_accuracy: 0.9632 - val\_loss: 0.1465


Epoch 65/100

**152/152**  **1s** 3ms/step - accuracy: 0.9920 - loss: 0.0210 - val\_accuracy: 0.9632 - val\_loss: 0.1463


Epoch 66/100

**152/152**  **1s** 3ms/step - accuracy: 0.9889 - loss: 0.0260 - val\_accuracy: 0.9632 - val\_loss: 0.1323


Epoch 67/100

**152/152**  **0s** 2ms/step - accuracy: 0.9932 - loss: 0.0243 - val\_accuracy: 0.9711 - val\_loss: 0.1196


Epoch 68/100

**152/152**  **1s** 2ms/step - accuracy: 0.9929 - loss: 0.0224 - val\_accuracy: 0.9658 - val\_loss: 0.1347


Epoch 69/100

**152/152**  **1s** 2ms/step - accuracy: 0.9959 - loss: 0.0150 - val\_accuracy: 0.9684 - val\_loss: 0.1464

Epoch 70/100

**152/152**  **0s** 2ms/step - accuracy: 0.9975 - loss: 0.0110 - val\_accuracy: 0.9684 - val\_loss: 0.1295


Epoch 71/100

**152/152**  **1s** 2ms/step - accuracy: 0.9884 - loss: 0.0240 - val\_accuracy: 0.9605 - val\_loss: 0.1780


Epoch 72/100

**152/152**  **0s** 2ms/step - accuracy: 0.9908 - loss: 0.0208 - val\_accuracy: 0.9711 - val\_loss: 0.1416

Epoch 73/100



















**152/152**  **0s** 2ms/step - accuracy: 0.9952 - loss: 0.0186 - val\_accuracy: 0.9579 - val\_loss: 0.2026

Epoch 74/100

**152/152**  **0s** 2ms/step - accuracy: 0.9934 - loss: 0.0197 - val\_accuracy: 0.9658 - val\_loss: 0.1514

Epoch 75/100


**152/152**  **1s** 2ms/step - accuracy: 0.9912 - loss: 0.0167 -

```
val_accuracy: 0.9684 - val_loss: 0.1610
Epoch 76/100
152/152  1s 2ms/step - accuracy: 0.9901 - loss: 0.0177 -
val_accuracy: 0.9553 - val_loss: 0.2135
Epoch 77/100
152/152  1s 2ms/step - accuracy: 0.9919 - loss: 0.0281 -
val_accuracy: 0.9658 - val_loss: 0.1849
Epoch 78/100
152/152  0s 2ms/step - accuracy: 0.9948 - loss: 0.0104 -
val_accuracy: 0.9658 - val_loss: 0.1889
Epoch 79/100
152/152  0s 2ms/step - accuracy: 0.9902 - loss: 0.0223 -
val_accuracy: 0.9605 - val_loss: 0.1822
Epoch 80/100
152/152  0s 2ms/step - accuracy: 0.9934 - loss: 0.0171 -
val_accuracy: 0.9605 - val_loss: 0.1741
Epoch 81/100
152/152  0s 2ms/step - accuracy: 0.9919 - loss: 0.0223 -
val_accuracy: 0.9605 - val_loss: 0.1972
Epoch 82/100
152/152  1s 2ms/step - accuracy: 0.9911 - loss: 0.0166 -
val_accuracy: 0.9658 - val_loss: 0.1552
Epoch 83/100
152/152  0s 2ms/step - accuracy: 0.9948 - loss: 0.0284 -
val_accuracy: 0.9632 - val_loss: 0.1833
Epoch 84/100
152/152  0s 2ms/step - accuracy: 0.9970 - loss: 0.0109 -
val_accuracy: 0.9632 - val_loss: 0.1796
Epoch 85/100
152/152  0s 2ms/step - accuracy: 0.9943 - loss: 0.0153 -
val_accuracy: 0.9632 - val_loss: 0.1616
Epoch 86/100
152/152  1s 2ms/step - accuracy: 0.9929 - loss: 0.0112 -
val_accuracy: 0.9632 - val_loss: 0.1835
Epoch 87/100
152/152  1s 3ms/step - accuracy: 0.9949 - loss: 0.0115 -
val_accuracy: 0.9658 - val_loss: 0.1814
Epoch 88/100
152/152  1s 3ms/step - accuracy: 0.9977 - loss: 0.0066 -
val_accuracy: 0.9632 - val_loss: 0.1634
Epoch 89/100
152/152  1s 4ms/step - accuracy: 0.9938 - loss: 0.0170 -
val_accuracy: 0.9658 - val_loss: 0.1594
Epoch 90/100
152/152  1s 3ms/step - accuracy: 0.9942 - loss: 0.0178 -
val_accuracy: 0.9579 - val_loss: 0.2024
Epoch 91/100
152/152  1s 3ms/step - accuracy: 0.9988 - loss: 0.0071 -
val_accuracy: 0.9632 - val_loss: 0.1959
Epoch 92/100
152/152  1s 3ms/step - accuracy: 0.9995 - loss: 0.0078 -
val_accuracy: 0.9632 - val_loss: 0.1973
Epoch 93/100
152/152  0s 2ms/step - accuracy: 0.9972 - loss: 0.0078 -
val_accuracy: 0.9632 - val_loss: 0.1646
Epoch 94/100
```

```

152/152 ————— 1s 2ms/step - accuracy: 0.9967 - loss: 0.0182 -
val_accuracy: 0.9632 - val_loss: 0.1945
Epoch 95/100
152/152 ————— 1s 2ms/step - accuracy: 0.9954 - loss: 0.0116 -
val_accuracy: 0.9632 - val_loss: 0.1934
Epoch 96/100
152/152 ————— 0s 2ms/step - accuracy: 0.9884 - loss: 0.0254 -
val_accuracy: 0.9632 - val_loss: 0.2229
Epoch 97/100
152/152 ————— 1s 4ms/step - accuracy: 0.9970 - loss: 0.0067 -
val_accuracy: 0.9632 - val_loss: 0.2225
Epoch 98/100
152/152 ————— 0s 2ms/step - accuracy: 0.9978 - loss: 0.0099 -
val_accuracy: 0.9632 - val_loss: 0.2026
Epoch 99/100
152/152 ————— 1s 3ms/step - accuracy: 0.9986 - loss: 0.0063 -
val_accuracy: 0.9658 - val_loss: 0.2071
Epoch 100/100
152/152 ————— 1s 2ms/step - accuracy: 0.9943 - loss: 0.0159 -
val_accuracy: 0.9658 - val_loss: 0.2095

```



```

In [53]: accuracy_3 = model_3.evaluate(X_test, y_test)
print(f"Accuracy: {accuracy_3[1]}")

```

```

15/15 ————— 0s 2ms/step - accuracy: 0.9669 - loss: 0.2060
Accuracy: 0.9684873819351196

```

## Comparison of the models



## Original model

The original model shows high accuracy on the test set (95.17%). The accuracy curve for training and validation stabilizes after about 30 epochs, indicating that the model reaches early convergence. The loss function does not decrease significantly after that point, suggesting that further training does not bring additional improvements. The loss value (0.3307) indicates a reasonably good fit of the model.

## Model 2

Model 2, with a more complex architecture and more regularization, also achieves high accuracy on the test set (94.30%). The difference between training and validation accuracy is small, indicating that the model generalizes well. However, the loss value (0.1830) is lower compared to the original model, suggesting that the new architecture improves the model fit. Despite this, the accuracy is slightly lower, which could be the result of a more complex architecture that does not provide a significant improvement.

## Model 3

Model 3 shows similar test set accuracy (94.27%), but the loss function is considerably higher (0.6577), indicating that the model has more difficulty fitting the data. Despite the deeper architecture and the use of Dropout, the model does not seem to benefit as much from the additional complexity. It is possible that there is a slight overfitting or that the optimizer is not working as well for this specific case.

## General comparison

The original model and model 2 perform very similarly in terms of accuracy, with a slight advantage in loss for model 2. This indicates that model 2 may be a slightly better choice in terms of fit. Model 3, while also achieving high accuracy, shows a higher loss value, suggesting that the deeper architecture is not necessarily beneficial in this context and may require further adjustments to the hyperparameters to improve its performance.

Overall, **model 2** appears to be the most balanced in terms of accuracy and loss.